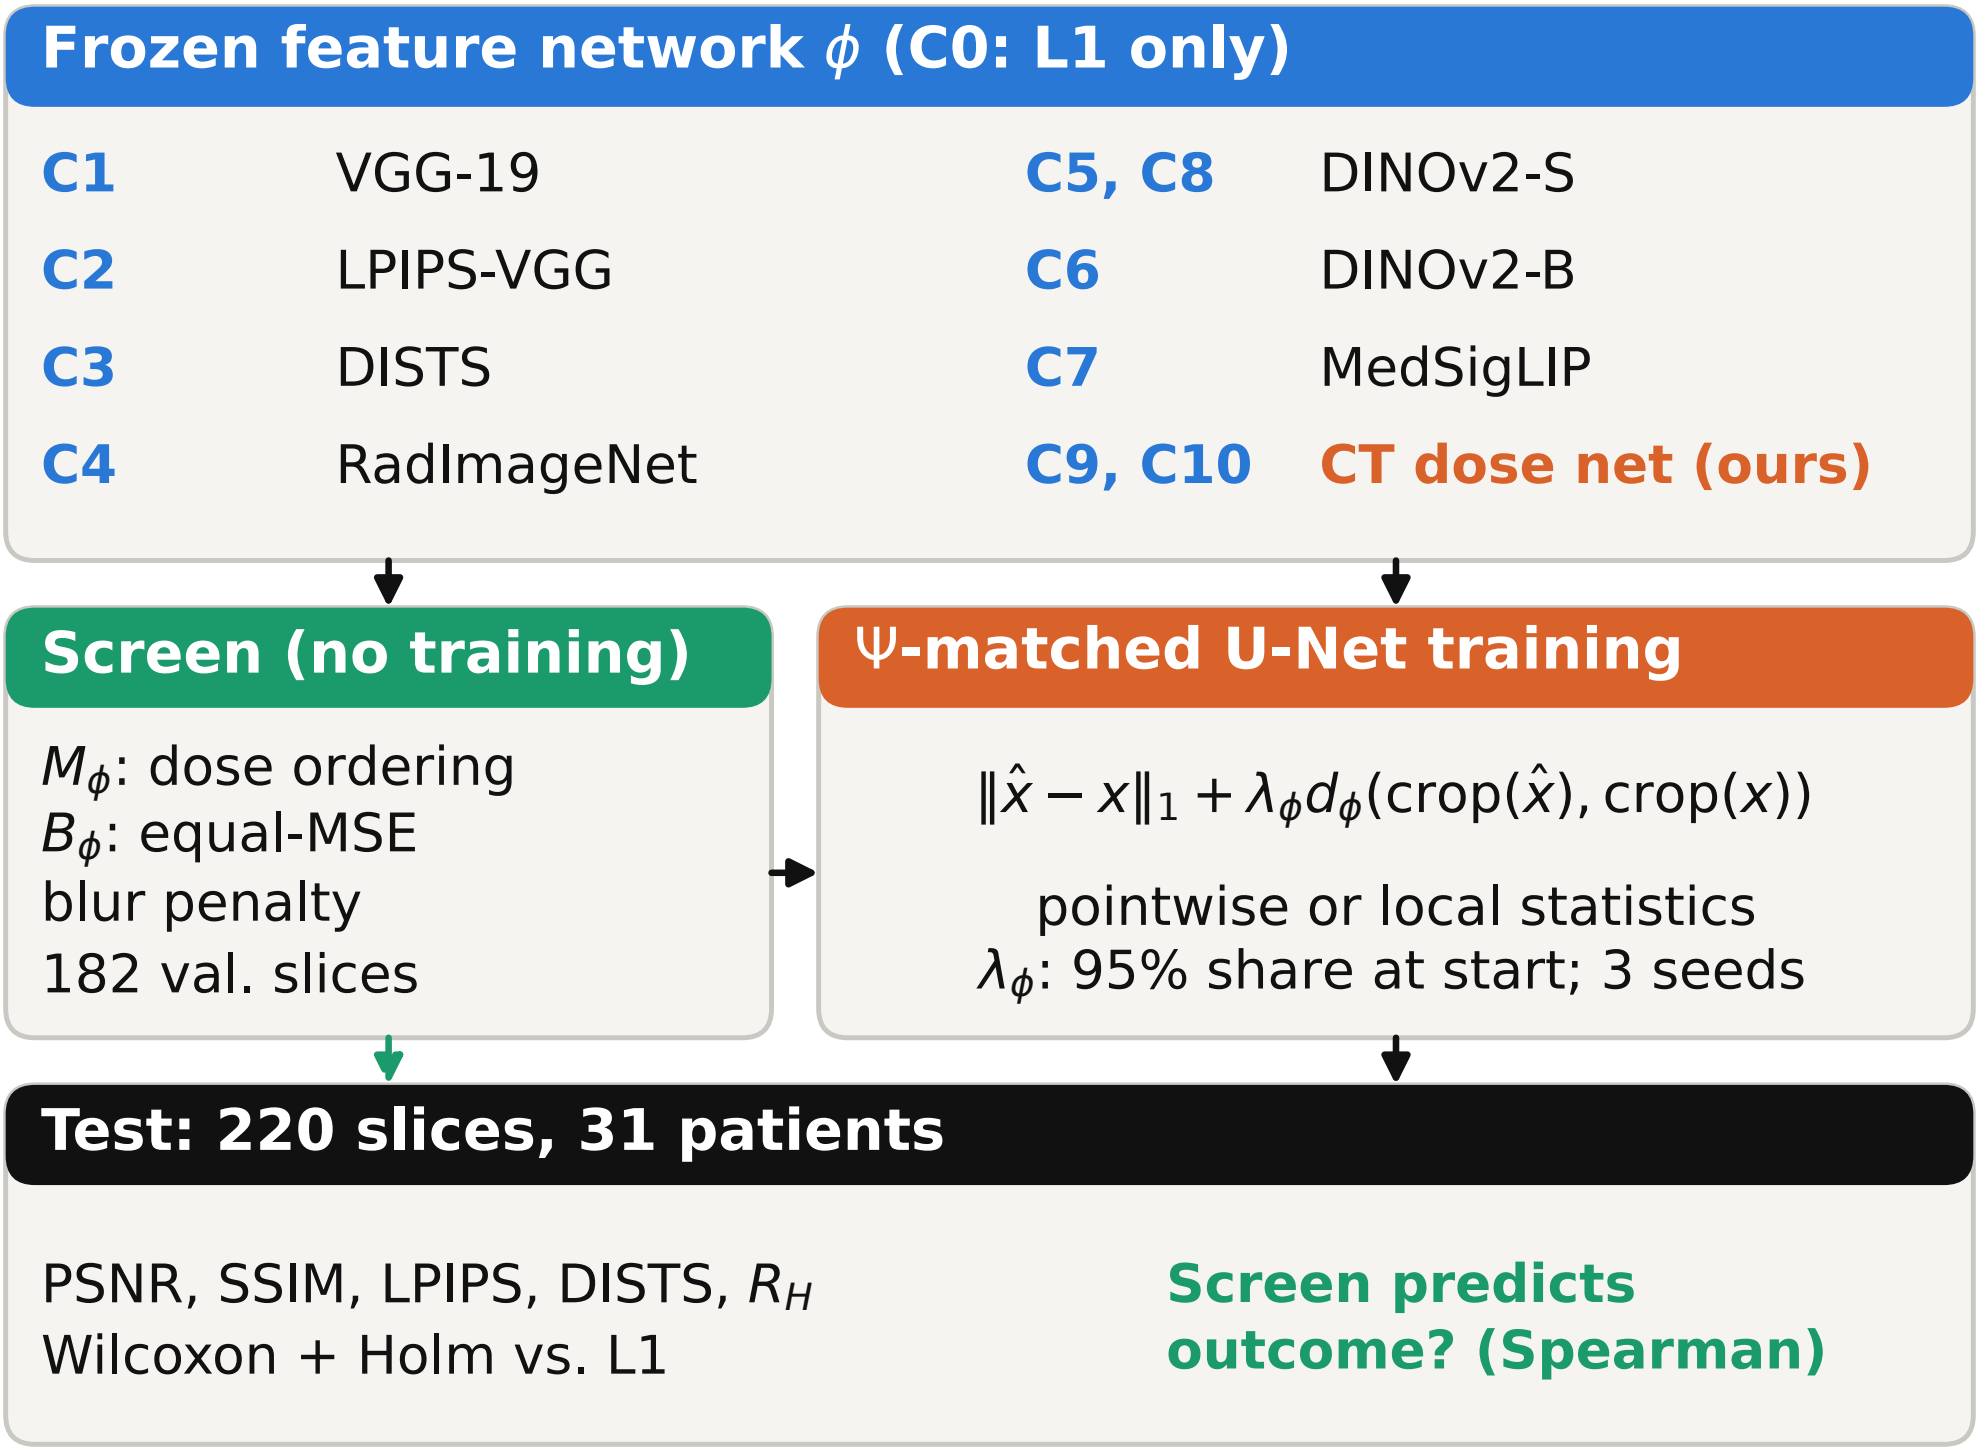
Frozen feature network ϕ (C0: L1 only)
C1 VGG-19
C2 LPIPS-VGG
C3 DISTS
C4 RadImageNet
C5, C8 DINOv2-S
C6 DINOv2-B
C7 MedSigLIP
C9, C10 CT dose net (ours)
Screen (no training)
M_ϕ: dose ordering
B_ϕ: equal-MSE
blur penalty
182 val. slices
Ψ-matched U-Net training
$\|\hat{x} - x\|_1 + \lambda_\phi d_\phi(\mathrm{crop}(\hat{x}), \mathrm{crop}(x))$
pointwise or local statistics
λ_ϕ: 95% share at start; 3 seeds
Test: 220 slices, 31 patients
PSNR, SSIM, LPIPS, DISTS, R_H
Wilcoxon + Holm vs. L1
Screen predicts outcome? (Spearman)